

CS701 - PPA - 1

Open Multiprocessing (OpenMP)

Team Members

Keerti Chaudhary (181CO226) Trivedi Naman
Manishbhai(181CO156)

1)

```
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q1$ g++ -o hello q1.cpp -fopenmp
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q1$ ./hello
Hello: 0
World: 0
Hello: 1
World: 1
Hello: 3
World: 3
Hello: 2
World: 2
Hello: 4
World: 4
```

The number of Threads = 5. Hence “Hello” and “World” are printed 5 times.

2)

```
ments/OpenMP/Q2$ ./a.out
Hello World! 3
Hello World! 2
Hello World! 0
Hello World! 1
```


The number of threads is 4. Hence thread ID ranges from 0 to 3.

3)

```
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q3$ g++ -o daxpy q3.cpp -fopenmp
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q3$ ./daxpy
Num_threads    Exec Time (micro-s)
1              534
2              752
3              270
4              318
5              506
6              600
7              597
8              2058
9              3278
10             343
11             344
12             352
13             330
14             320
15             329
16             334
17             351
18             392
19             325
20             315
```

When the number of threads is increased further, runtime increases since the number of threads are greater than the number of cores in the system. Hence time is consumed in scheduling the excess threads by operating system among the cores.

Exec time vs no of threads




Speedup = T_1/T_n , hence maximum speedup will be when T_n is smaller, therefore maximum speedup will be when the number of threads = 3

4)

```
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q4$ g++ -o matrix q4.cpp -fopenmp
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q4$ ./matrix
Num_threads    Exec Time (micro-s)
1              57881
2              31154
3              21797
4              18183
5              26631
6              22963
7              20790
8              18631
9              22161
10             22762
11             21175
12             22022
13             21504
14             19899
15             19327
16             18505
17             20743
18             19598
19             20181
20             20322
21             20164
22             19467
23             19235
24             18956
25             19902
26             19952
27             20538
```

Matrix multiplication done using parallel calculations speeds up the process as seen in the output. The matrix multiplication algorithm 3 loops nested with a complexity of $O(n^3)$. We fork the master thread at each level of the loop to obtain the best speedup. The innermost loop sums up the product of the corresponding row and column. This is done using the reduction operation provided by OpenMP.



5)

```
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q5$ ./calculate
Value of Pi by Sequential Calculation : 3.141593
Parallel Calculation
Pi : 3.141593    Speedup: 0.489552    Threads : 2
Pi : 3.141593    Speedup: 0.332804    Threads : 3
Pi : 3.141593    Speedup: 0.271645    Threads : 4
Pi : 3.141593    Speedup: 0.239640    Threads : 5
Pi : 3.141593    Speedup: 0.207211    Threads : 6
Pi : 3.141593    Speedup: 0.181080    Threads : 7
Pi : 3.141593    Speedup: 0.221869    Threads : 8
Pi : 3.141593    Speedup: 0.255792    Threads : 9
Pi : 3.141593    Speedup: 0.232992    Threads : 10
Pi : 3.141593    Speedup: 0.214680    Threads : 11
Pi : 3.141593    Speedup: 0.206145    Threads : 12
Pi : 3.141593    Speedup: 0.185534    Threads : 13
Pi : 3.141593    Speedup: 0.192922    Threads : 14
Pi : 3.141593    Speedup: 0.167830    Threads : 15
Pi : 3.141593    Speedup: 0.186283    Threads : 16
Pi : 3.141593    Speedup: 0.202870    Threads : 17
Pi : 3.141593    Speedup: 0.200251    Threads : 18
Pi : 3.141593    Speedup: 0.197134    Threads : 19
Pi : 3.141593    Speedup: 0.182535    Threads : 20
```

Shared variables are

1. sum: which stores the area under the curve

Private variables are

1. ID: thread ID
2. increment: loop increment to split the loop iteration among the threads
3. t_sum: to calculate the sum for each thread

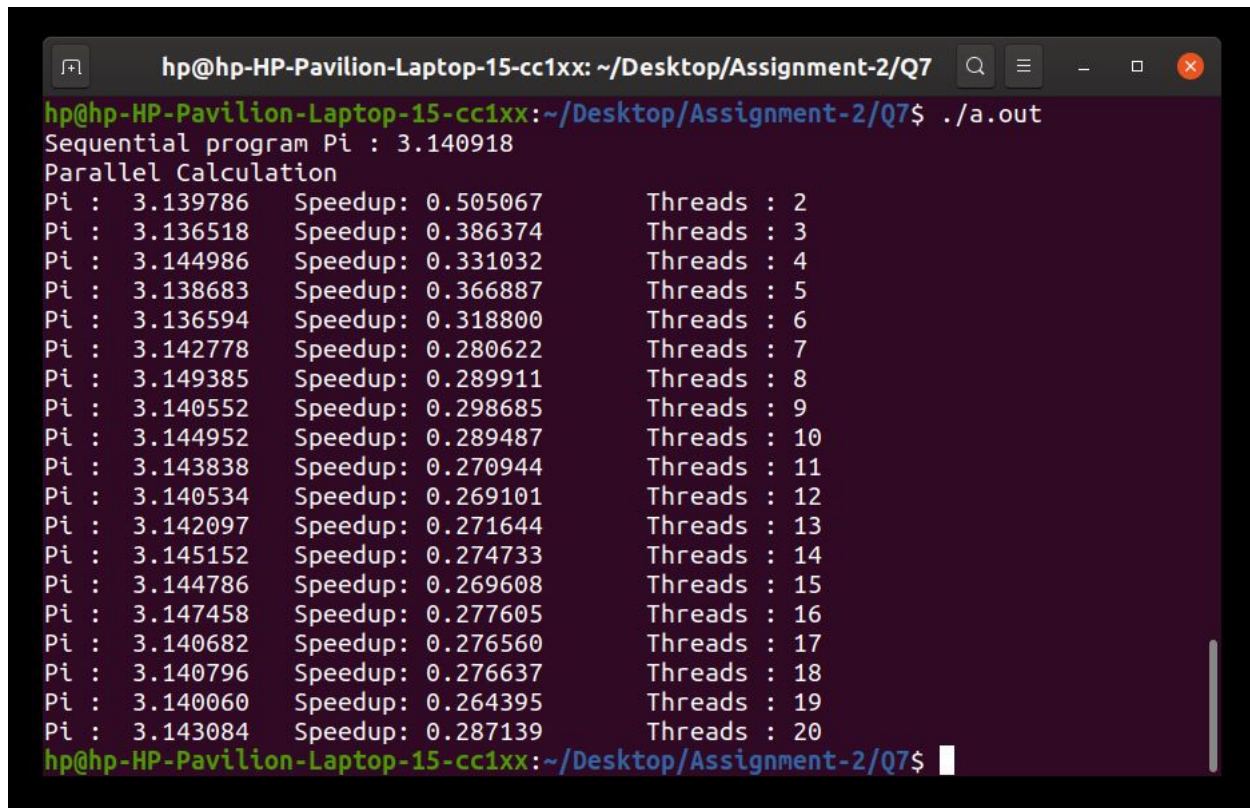
Since every thread spawned calculates its sum, the total sum has to be summed up. This sum variable is shared but cannot be simultaneously updated by multiple threads. Hence its updating is given inside a “critical” construct which ensures that only one thread updates its value at a time.

6)

```
hp@hp-HP-Pavilion-Laptop-15-cc1xx: ~/Desktop/Assignment-2/Q6
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q6$ ./calculate
Value of Pi by Sequential Calculation : 3.141593
Time taken : 9.629138 ms
Parallel Calculation
Pi : 3.141593 Speedup: 0.455915 Threads : 2
Pi : 3.141593 Speedup: 0.314556 Threads : 3
Pi : 3.141593 Speedup: 0.251994 Threads : 4
Pi : 3.141593 Speedup: 0.210131 Threads : 5
Pi : 3.141593 Speedup: 0.192873 Threads : 6
Pi : 3.141593 Speedup: 0.162825 Threads : 7
Pi : 3.141593 Speedup: 0.167198 Threads : 8
Pi : 3.141593 Speedup: 0.228843 Threads : 9
Pi : 3.141593 Speedup: 0.210472 Threads : 10
Pi : 3.141593 Speedup: 0.207313 Threads : 11
Pi : 3.141593 Speedup: 0.183719 Threads : 12
Pi : 3.141593 Speedup: 0.171179 Threads : 13
Pi : 3.141593 Speedup: 0.164727 Threads : 14
Pi : 3.141593 Speedup: 0.171988 Threads : 15
Pi : 3.141593 Speedup: 0.171509 Threads : 16
Pi : 3.141593 Speedup: 0.191728 Threads : 17
Pi : 3.141593 Speedup: 0.187165 Threads : 18
Pi : 3.141593 Speedup: 0.174548 Threads : 19
Pi : 3.141593 Speedup: 0.171058 Threads : 20
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q6$
```

The reduction clause is used to sum the area under the curve. Every thread creates its local copy of the sum and then the reduction is done using the operator specified in *reduction (op: list)*.

7)



```
hp@hp-HP-Pavilion-Laptop-15-cc1xx: ~/Desktop/Assignment-2/Q7$ ./a.out
Sequential program Pi : 3.140918
Parallel Calculation
Pi : 3.139786 Speedup: 0.505067 Threads : 2
Pi : 3.136518 Speedup: 0.386374 Threads : 3
Pi : 3.144986 Speedup: 0.331032 Threads : 4
Pi : 3.138683 Speedup: 0.366887 Threads : 5
Pi : 3.136594 Speedup: 0.318800 Threads : 6
Pi : 3.142778 Speedup: 0.280622 Threads : 7
Pi : 3.149385 Speedup: 0.289911 Threads : 8
Pi : 3.140552 Speedup: 0.298685 Threads : 9
Pi : 3.144952 Speedup: 0.289487 Threads : 10
Pi : 3.143838 Speedup: 0.270944 Threads : 11
Pi : 3.140534 Speedup: 0.269101 Threads : 12
Pi : 3.142097 Speedup: 0.271644 Threads : 13
Pi : 3.145152 Speedup: 0.274733 Threads : 14
Pi : 3.144786 Speedup: 0.269608 Threads : 15
Pi : 3.147458 Speedup: 0.277605 Threads : 16
Pi : 3.140682 Speedup: 0.276560 Threads : 17
Pi : 3.140796 Speedup: 0.276637 Threads : 18
Pi : 3.140060 Speedup: 0.264395 Threads : 19
Pi : 3.143084 Speedup: 0.287139 Threads : 20
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q7$
```

Random numbers are generated using the **Linear Congruential Generator** defined by

$$X(n+1) = (a \cdot X(n) + c) \bmod m$$

Bonus part

1. Random number generator was made threadsafe using the “#pragma omp thread private(seed)” construct in OpenMP. This allows each thread to have its private seed.
2. The random number generator was made numerically correct, ie with non-overlapping sequences of pseudo-random numbers by using the leapfrog algorithm.

This can be done for linear congruential generators and involves replacing the multiplier a and the additive constant c by new values a^N and $c(a^N - 1)/(a - 1)$, where N is the number of threads.

8)

```
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q8$ gcc -o pc q8.c -fopenmp
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q8$ ./pc
Producer populated data
Consumer calculated Array sum
In 0.002724 seconds, Sequential code gives sum : 100000.000000
Producer populated data
Consumer calculated Array sum
In 0.000460 seconds, Parallel code gives sum : 100000.000000
Speed up : 0.168758
```

Two threads, one for producer and one for the consumer was created. The sections clause was used to reserve each section for one thread.

The consumer knows the producer has completed execution by getting notified by a global *flag*.

9)

```
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q9$ gcc -o pc q9.c -fopenmp
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q9$ ./pc
Buffer empty
Thread 1 Produced A
Thread 0 Consumed A
Thread 1 Produced B
Thread 1 Produced C
Thread 0 Consumed B
Thread 0 Consumed C
Thread 1 Produced D
Thread 1 Produced E
Thread 0 Consumed D
Thread 1 Produced F
Thread 0 Consumed E
Thread 0 Consumed F
Thread 1 Produced G
Thread 0 Consumed G
Thread 1 Produced H
Thread 1 Produced I
Thread 0 Consumed H
Thread 1 Produced J
Thread 0 Consumed I
```

Here we have to introduce the sleep statement to make sure that that none of them are running multiple times without the other getting a chance.

As for the race condition, the critical section problem may occur due to both updating the same values. So we kept it under the critical section. OpenMP gives APIs who implement locks for us to ensure the atomic behavior for that section.

10)

```
@naman-Lenovo-ideapad-530S-15IKB: ~/Downloads/M5-1-OpenMP/md-example
Thread 3 Produced W
Thread 0 Consumed W
Thread 2 Consumed W
Thread 1 Produced X
Thread 3 Produced X
Thread 0 Consumed X
Thread 2 Consumed X
naman@naman-Lenovo-ideapad-530S-15IKB:~/Downloads/M5-1-OpenMP/md-example$ gcc prod1.c -fopenmp
naman@naman-Lenovo-ideapad-530S-15IKB:~/Downloads/M5-1-OpenMP/md-example$ ./a.out
Thread 1 Produced A
Thread 2 Consumed A
Buffer empty
Thread 3 Produced A
Thread 1 Produced B
Thread 2 Consumed A
Thread 0 Consumed B
Thread 3 Produced B
Thread 2 Consumed B
Buffer empty
Thread 3 Produced C
Thread 1 Produced C
Thread 0 Consumed C
Thread 2 Consumed C
Thread 3 Produced D
Thread 1 Produced D
Thread 2 Consumed D
Thread 3 Produced E
Thread 0 Consumed D
Thread 1 Produced E
Thread 0 Consumed E
Thread 3 Produced F
Thread 1 Produced F
Thread 2 Consumed E
^C
```

Here to avoid a producer or consumer entering while any other producer or consumer is in the critical section we define a global critical section. Synchronous behaviour can not be achieved as there may be too many processes in the ready queue of some a processor executing a particular thread thereby delaying its turn. Here using a critical section helps avoids us from ending into an infinite loop.

11)

```
hp@hp-HP-Pavilion-Laptop-15-cc1xx: ~/Desktop/Assignment-2/Q11
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q11$ gcc -o md q11.c -lm -fopenmp
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q11$ ./md
Molecular Dynamics Simulation example program
-----
number of particles is ..... 13500
side length of the box is ..... 25.323179
cut off is ..... 3.750000
reduced temperature is ..... 0.722000
basic timestep is ..... 0.064000
temperature scale interval ..... 10
stop scaling at move ..... 20
print interval ..... 5
total no. of steps ..... 20


  t      ke      pe      e      temp      pres      vel      rp
  ---      -      -      -      -      -      -      -
   5 12619.1898 -91985.4707 -79366.2809  0.6232 -5.2874  0.1821 39.7
  10 14619.4170 -86198.9064 -71579.4894  0.7220 -2.8280  0.1339 14.4
  15 11335.9135 -82913.7738 -71577.8603  0.5598 -1.4860  0.1711 33.6
  20 10796.2210 -82374.2671 -71578.0461  0.5332 -1.2110  0.1676 32.1
Time = 19.059803
hp@hp-HP-Pavilion-Laptop-15-cc1xx:~/Desktop/Assignment-2/Q11$
```

The key things to note in the parallelized code:-

- Here loop iterations are independent so the order of their execution doesn't matter and thus they can be parallelized
- Here we have parallelized the outer food loop
- Variable j needs to be private as every loop might be executing different iterations
- We don't need to bother with the local variables of the section
- Variables epot and vir are being accumulated in every iteration hence they will be our reduction variables
- There are sections in which $f[i]$ and $f[j]$ are updated, we do not need to worry about $f[i]$ as it has been taken care of by omp for, but the $f[j]$ update needs to be handled as it should not be updated simultaneously by all the threads, values j might be same for many, thus it is being declared under the critical section.

Ans1. How does the data-sharing change?

This is the general functioning of omp. The task is distributed equally for iterations amongst all the threads and in the end, partial-sums of each thread are summed together to get the final result.



Ans2. Here as there was a race condition in the update of `f[j]` thus we had to introduce `omp atomic/critical` for resolving it. So if we use different arrays for each thread then there won't occur any race condition, this will save us from bottleneck occurring due to atomic operations. In the end, they are going to be executed serially after this for a loop as one thread, the master thread thus the value in the master thread should be correct and accumulation of all.